

IERG 3050 Assignment 3

Due: 11 April 2025

- Submit a single .pdf file containing all your answers to the Blackboard before the due date.
- Answer all questions.
- Type or write your work neatly.

1. We want to predict the weight (y) of whitefishes from their heights (h) and widths (w). Here are the data obtained from a fish market:

y	h	w
270	8.3804	4.2476
270	8.1454	4.2485
306	8.778	4.6816
540	10.744	6.562
800	11.7612	6.5736
1000	12.354	6.525

Find the plane $y = a_0 + a_1h + a_2w$ by multiple linear regression. What is the predicted weight of a whitefish of height 10 and width 5.5?

2. Apply Monte Carlo simulation with importance sampling to compute $\int_0^2 \frac{dx}{\sqrt{x}}$.
3. Generate 300 samples of $U(0,1)$ by any programming language or software you like. Perform the uniformity test of these samples by chi-square goodness-of-fit test with 5 equal-length bins and at 5% significance level. Also show the histogram of the bins.
4. Generate 5 samples from a Student's t -distribution with 30 degrees of freedom, and then standardize them. Sort the values in ascending order, and let the values be the 0.1-, 0.3-, 0.5-, 0.7- and 0.9-quantiles. Plot the Q-Q plot of these values against the standard normal distribution. What is the conclusion of the Q-Q plot?

1. We want to predict the weight (y) of whitefishes from their heights (h) and widths (w). Here are the data obtained from a fish market:

y	h	w
270	8.3804	4.2476
270	8.1454	4.2485
306	8.778	4.6816
540	10.744	6.562
800	11.7612	6.5736
1000	12.354	6.525

Find the plane $y = a_0 + a_1h + a_2w$ by multiple linear regression. What is the predicted weight of a whitefish of height 10 and width 5.5?

```

q1.py > ...
1  import pandas as pd
2
3  data = {
4      'y': [270, 270, 306, 540, 800, 1000],
5      'h': [8.3804, 8.1454, 8.778, 10.744, 11.7612, 12.354],
6      'w': [4.2476, 4.2485, 4.6816, 6.562, 6.5736, 6.525]
7  }
8  df = pd.DataFrame(data)
9  from sklearn.linear_model import LinearRegression
10
11  X = df[['h', 'w']]
12  Y = df['y']
13
14  regr = LinearRegression()
15  regr.fit(X, Y)
16
17  print('Intercept:', regr.intercept_)
18  print('Coefficients:', regr.coef_)
19
20  new_h = 10
21  new_w = 5.5
22  predicted_y = regr.predict([[new_h, new_w]])
23  print('Predicted Weight:', predicted_y[0])

```

```

(.venv) lifehater@LifedeMacBook-Pro temp %
win-arm64/bundled/libs/debugpy/adapters/...
Intercept: -1263.789283808669
Coefficients: [ 275.77358715 -177.31218183]
Predicted Weight: 518.7295875821146

```

2. Apply Monte Carlo simulation with importance sampling to compute $\int_0^2 \frac{dx}{\sqrt{x}}$.

For $I = \int_0^2 \frac{1}{\sqrt{x}} dx$

Exact analytical solution:

$$I = \int_0^2 x^{-\frac{1}{2}} dx = \left[2x^{\frac{1}{2}} \right]_0^2 = 2\sqrt{2}$$

For $f(x) = \frac{1}{\sqrt{x}} \Rightarrow$ we select $p(x) = \frac{1}{2\sqrt{x}}$ for $x \in (0, 2]$

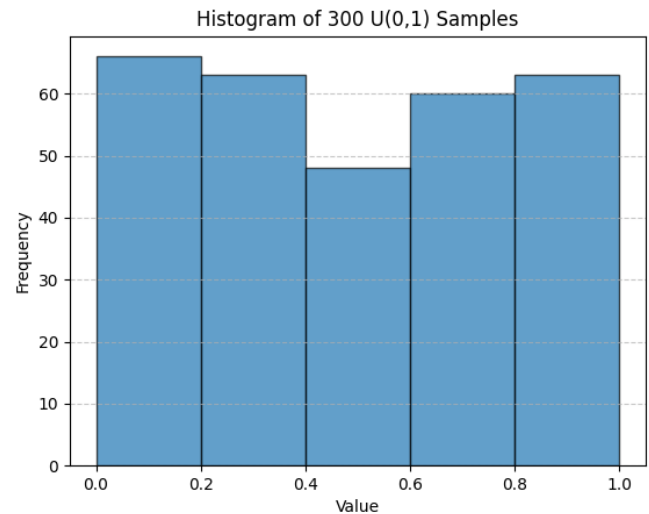
Importance sampling estimate: $I \approx \frac{1}{N} \sum_{i=1}^N \frac{f(x_i)}{p(x_i)} = \frac{1}{N} \sum_{i=1}^N \frac{\frac{1}{\sqrt{x_i}}}{\frac{1}{2\sqrt{x_i}}} = \frac{1}{N} \sum_{i=1}^N 2 = 2 \quad (\times)$

Correct normalized distribution: $p(x) = \frac{1}{2\sqrt{2}x}$

$$I \approx \frac{1}{N} \sum_{i=1}^N \frac{\frac{1}{\sqrt{x_i}}}{\frac{1}{2\sqrt{2}\sqrt{x_i}}} = \frac{1}{N} \sum_{i=1}^N 2\sqrt{2} = 2\sqrt{2} \quad //$$

3. Generate 300 samples of $U(0,1)$ by any programming language or software you like. Perform the uniformity test of these samples by chi-square goodness-of-fit test with 5 equal-length bins and at 5% significance level. Also show the histogram of the bins.

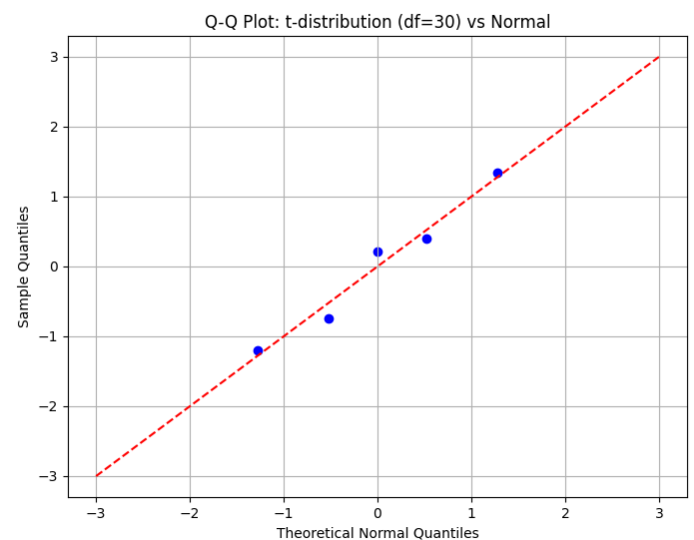
```
q3.py > ...
1 import numpy as np
2 import matplotlib.pyplot as plt
3 from scipy.stats import chisquare
4
5 # 1. Generate 300 samples from U(0, 1)
6 np.random.seed(42) # For reproducibility
7 samples = np.random.uniform(0, 1, 300)
8
9 # 2. Define 5 equal-length bins between 0 and 1
10 bin_edges = np.linspace(0, 1, 6) # 6 edges for 5 bins
11 observed_counts, _ = np.histogram(samples, bins=bin_edges)
12
13 # 3. Expected counts for each bin (uniform): 300 / 5 = 60
14 expected_counts = np.array([60] * 5)
15
16 # 4. Perform chi-square goodness-of-fit test
17 chi2_stat, p_value = chisquare(f_obs=observed_counts, f_exp=expected_counts)
18
19 # 5. Output results
20 print("Observed counts:", observed_counts)
21 print("Expected counts:", expected_counts)
22 print(f"Chi-square statistic: {chi2_stat:.4f}")
23 print(f"P-value: {p_value:.4f}")
24
25 if p_value < 0.05:
26     print("Reject the null hypothesis: Data is not uniformly distributed.")
27 else:
28     print("Fail to reject the null hypothesis: Data may be uniformly distributed.")
29
30 # 6. Plot histogram
31 plt.hist(samples, bins=bin_edges, edgecolor='black', alpha=0.7)
32 plt.title("Histogram of 300 U(0,1) Samples")
33 plt.xlabel("Value")
34 plt.ylabel("Frequency")
35 plt.xticks(bin_edges)
36 plt.grid(axis='y', linestyle='--', alpha=0.7)
37 plt.show()
```



```
(.venv) lifehater@LifedeMacBook-Pro temp %
win-arm64/bundled/libs/debugpy/adapter/../../
Observed counts: [66 63 48 60 63]
Expected counts: [60 60 60 60 60]
Chi-square statistic: 3.3000
P-value: 0.5089
```

4. Generate 5 samples from a Student's t -distribution with 30 degrees of freedom, and then standardize them. Sort the values in ascending order, and let the values be the 0.1-, 0.3-, 0.5-, 0.7- and 0.9-quantiles. Plot the Q-Q plot of these values against the standard normal distribution. What is the conclusion of the Q-Q plot?

```
q4.py > ...
1 import numpy as np
2 import scipy.stats as stats
3 import matplotlib.pyplot as plt
4
5 # Set random seed for reproducibility
6 np.random.seed(42)
7
8 # 1. Generate 5 samples from t-distribution with 30 df
9 t_samples = stats.t.rvs(df=30, size=5)
10
11 # 2. Standardize the samples (though t with high df is already ~N(0,1))
12 standardized = (t_samples - np.mean(t_samples)) / np.std(t_samples, ddof=1)
13
14 # 3. Sort the values
15 sorted_samples = np.sort(standardized)
16
17 # 4. Define desired quantiles
18 quantiles = np.array([0.1, 0.3, 0.5, 0.7, 0.9])
19
20 # 5. Get theoretical normal quantiles
21 normal_quantiles = stats.norm.ppf(quantiles)
22
23 # 6. Create Q-Q plot
24 plt.figure(figsize=(8, 6))
25 plt.scatter(normal_quantiles, sorted_samples, color='blue')
26 plt.plot([-3, 3], [-3, 3], 'r--') # y=x reference line
27 plt.title('Q-Q Plot: t-distribution (df=30) vs Normal')
28 plt.xlabel('Theoretical Normal Quantiles')
29 plt.ylabel('Sample Quantiles')
30 plt.grid(True)
31 plt.show()
32
33 # 7. Calculate correlation for additional insight
34 correlation = np.corrcoef(normal_quantiles, sorted_samples)[0, 1]
35 print(f"Correlation between theoretical and sample quantiles: {correlation:.4f}")
```



```
Correlation between theoretical and sample quantiles: 0.9850
```